

SIO221a_Lecture10_MATLAB

October 31, 2023

SIO221a Notes - Alford and Gille

Reading: Bendat and Piersol Ch. 4.2.2

Concepts covered: Spectra and their uncertainties; Parseval's theorem for spectra; Fourier transform of a boxcar

1 Lecture 10

Recap So far this class has looked at methods for computing spectra. We made sure that we understood how to normalize our amplitudes, we segmented data so that we'd have something to average, and we detrended to remove the bias associated with large-scale trends. We've also identified the the lowest resolved frequency and the Nyquist frequency (the highest frequency for which we have any information.) We've almost managed to produce useful spectra, but we need to assign error bars and figure out how many degrees of freedom we have, and we need to make sure that everything is normalized correctly and add units. We'll also need to put a little consideration into edge effects.

Spectral Uncertainties The problem with all of the spectra that we've computed so far (the amplitude of the Fourier Transform) is that we have no way to evaluate the uncertainty. By eye we can see that it's fairly noisy. We know from computing means that the error of the mean decreases as we average more quantities together.

How do we incorporate more data into our spectra? You might imagine that you could improve your spectrum by extending the input time series from N to $2N$ data points for example. Unfortunately, although adding data points will change your spectrum, it won't reduce your noise or make the spectrum more precise at any individual frequency. Instead it will increase the number of frequencies for which you obtain results from $N/2$ to N .

Error bars for spectra rely on a principle similar to what we used when we estimated the standard error of the mean. Our uncertainties in our spectra decrease as we average more spectra together. The challenge is to figure out how to obtain more spectra that can be averaged together. Typically what we do is to break our time series into segments, compute spectra for each of the segments, and average these to get a mean spectrum. Since we're averaging squared quantities (the spectral amplitudes) the distribution of the sum will follow a χ^2 distribution. And we'll use this to estimate the uncertainties. When we compute our estimate of the spectrum of our data, we'll compute coefficients $a(f)$, and for n segments, at frequency f , the estimated spectrum will be

$$\hat{E}(f) = \langle |a(f)|^2 \rangle = \frac{1}{n} \sum_{i=1}^n |a(f)|^2$$

(Since we actually had two degrees of freedom for each segment, this will give us $\nu = 2n$ degrees of freedom.)

Now to estimate errors the whole argument stems from a notion that there is a true spectrum $E(f)$, which we don't know, and our best estimate $\hat{E}(f)$, that we've computed, but which might be noisy. We could consider the difference between E and \hat{E} , but this gets hairy quickly. Besides we're going to look at our results in log/log space anyway. So let's look at the ratio of the two. Recall that $\ln(\hat{E}/E) = \ln(\hat{E}) - \ln(E)$. We can define a variable that represents this sum of squares normalized by the true spectrum

$$y = \nu \frac{\hat{E}_\nu(f)}{E(f)}$$

where \hat{E}_ν is our best estimate of the spectrum from real observations, when we have ν degrees of freedom.

The spectrum is an average (or sum) of squared amplitudes, which means that the ratio of \hat{E} to $E(f)$ should have a χ^2 distribution with ν degrees of freedom. Thus we'll always consider the ratio $\hat{E}(f)/E(f)$, where we use $\nu/2$ data segments.

Even though we don't know E , we know that the ratio of \hat{E} to E (\hat{E}/E) should be about 1. The spectrum is an average (or sum) of squared amplitudes, which makes it a χ^2 quantity, and since χ^2 statistics have an expectation value corresponding to the number of degrees of freedom, it turns out to be convenient to scale it by ν , so that y has an expectation value of ν . The standard deviation is $\langle y^2 \rangle - \langle y \rangle^2 = 2\nu$ (see Bendat and Piersol, section 4.2.2). From that we can infer that with two degrees of freedom, the standard deviation $\langle \hat{E}^2 \rangle - \langle \hat{E} \rangle^2 = 2/\nu E_0^2 = E_0^2$, which means that the standard deviation is equal to the original value. Clearly we need more samples.

Now in reality, we don't care about the details of χ^2 but rather the probability that our estimate of the spectrum \hat{E} is close to or far from the unknown true spectrum E , so we're going to look at the probability that $\nu \hat{E}(f)/E(f)$ falls within a fixed range. So we want to use the χ^2 distribution as a probability distribution and ask about the probability that $\hat{E}(f)/E(f)$ should be within \pm some range of 1. Formally, the probability that the estimated spectrum should be close in value to the true spectrum is:

$$P \left(\chi_{\nu, 1-\alpha/2}^2 < \nu \frac{\hat{E}(f)}{E(f)} < \chi_{\nu, \alpha/2}^2 \right) = 1 - \alpha$$

so if we want to find a 95% significance level, we set α to 0.05. We'll want to know the value of χ^2 that corresponds to a given point on the cdf, so for this we use the inverse χ^2 function ("chi2inv" in Matlab). So we can invert this relationship to provide the probability that the true value is within a particular range of the observed value:

$$P \left(\frac{\chi_{\nu, 1-\alpha/2}^2}{\nu} < \frac{\hat{E}(f)}{E(f)} < \frac{\chi_{\nu, \alpha/2}^2}{\nu} \right) = 1 - \alpha$$

$$P \left(\frac{\nu}{\chi_{\nu, 1-\alpha/2}^2} > \frac{E(f)}{\hat{E}(f)} > \frac{\nu}{\chi_{\nu, \alpha/2}^2} \right) = 1 - \alpha$$

$$P\left(\frac{\nu \hat{E}(f)}{\chi_{\nu, 1-\alpha/2}^2} > E(f) > \frac{\nu \hat{E}(f)}{\chi_{\nu, \alpha/2}^2}\right) = 1 - \alpha$$

$$P\left(\frac{\nu}{\chi_{\nu, \alpha/2}^2} < \frac{E(f)}{\hat{E}(f)} < \frac{\nu}{\chi_{\nu, 1-\alpha/2}^2}\right) = 1 - \alpha$$

Whether you use χ^2 or its reciprocal, the ratio between the high and low error bars should be the same, and the ratio will be what matters.

This error formulation differs from the usual error bars that we're used to seeing where we say for example that the true temperature should be the measured temperature plus or minus an uncertainty: $T = \hat{T} \pm \delta_T$. We can develop a similar expression for the true spectrum: $E(f)$ is in the range between $\nu \hat{E}(f) / \chi_{\nu, \alpha/2}^2$ and $\nu \hat{E}(f) / \chi_{\nu, 1-\alpha/2}^2$, where ν is twice the number of segments. This expression isn't very easy to interpret, since it varies as a function of frequency, and the estimated value $\hat{E}(f)$ is not at the mid-point of the range.

Instead we'll keep in mind that we've computed the uncertainty for the ratio $\hat{E}(f)/E(f)$, and the probabilities for this ratio does not depend on frequency. On a log plot, error bars defined by the range between $\nu/\chi_{\nu, \alpha/2}^2$ and $\nu/\chi_{\nu, 1-\alpha/2}^2$ are the same size at all frequencies, so we can easily compare spectral peaks at different frequencies.

Some statistics books include look-up tables for χ^2 , but we can compute it directly in Matlab. For N data segments and $\nu = 2 * N$ degrees of freedom, the error limits are:

```
[2]: nu=10;
    err_low = nu/chi2inv(1-.05/2,nu)
    err_high = nu/chi2inv(.05/2,nu)
```

```
err_high =
```

```
3.0798
```

```
err_low =
```

```
0.4882
```

We can plot these values as:

```
semilogy([f],[err_low err_high]*A);
```

where we set the frequency f and the amplitude A , so that the error bar ends up positioned in a convenient spot on the plot.

Now to have M data segments (and $\nu = 2M$ degrees of freedom), we have to split our long data record into shorter segments. We can do this by taking N data points at a time.

Let's first re-generate our spectrum from the last lecture:

```

[19]: %This would be a nice bit of code to turn into a function!

%Nominal example:
%First I'll need to pull the new function to my local working directory.
%[f,P]=Class_Spectrum_MHA(data,time,M);

%loglog(f,P);
%grid

%Let's load in some real data from Alford et al (2012) and have a look.
clear all
cd('/Users/malford/GoogleDrive/Work/Projects/Teaching/sio221a/MHA_2023/')
addpath('code')
load('data/Vel_2008-2009.mat')

iz=min(find(Vel.z > 300));

it=1:16000;
data=Vel.u(iz,it);
time=Vel.dtnum(it);
dt=nanmean(diff(time));

%Next, detrend and remove the mean.
data=detrend(data);

%
a=fft(data);
N=length(data);
amp=abs(a(1:N/2+1)).^2; % for even N
% amp=abs(a(1:(N+1)/2)).^2; % for odd N

%% which freqs do these correspond to?
T=dt*N;

df=1/T;
fn=1/2/dt;

f=0:df:fn; %frequency vector, cpd, goes from 0 to Nyquist.

amp = amp / N.^2; % first correct for the MATLAB normalization
amp = amp .* 2; %we threw out half of the spectrum; so correct for the lost
    ↪ variance.
amp = amp / df; % this is then the definition of the spectrum

M=8; p=N/M;
datab=reshape(data,N/M,M); % this gives us an array with N/M points
    % per column and M columns

```

```

b=fft(datab); % this computes the fft for each column
amp_b=abs(b(1:p/2+1,:)).^2;

%% Compute frequencies - note differences!
Tb=dt*p;

dfb=1/Tb;
fn=1/2/dt;

fb=0:dfb:fn; %frequency vector, cpd, goes from 0 to Nyquist.


%% Normalize as above
amp_b = amp_b / p.^2; % first correct for the MATLAB normalization
amp_b = amp_b .* 2; %we threw out half of the spectrum; so correct for the lost
    ↪ variance.
amp_b = amp_b / dfb; % this is then the definition of the spectrum

myspec=nanmean(amp_b,2);

%% Now compare - and let's zoom in a bit
loglog(f,amp,fb,myspec);
ylim([1e-7 1e-1])
xlim([1e-2 50])
grid

%% label
ylabel('\Phi_{Vel} / (m/2)^2 / cpd')
xlabel('frequency / cpd')

```



```
[24]: nu=2*M;
err_high = nu/chi2inv(.05/2,nu);
err_low = nu/chi2inv(1-.05/2,nu);

%semilogy(0:N/2,sd,[N/4 N/4],[err_low err_high]*sd(N/4))
indf=min(find(fb > 12))
myspec(indf)
err_high
h=loglog(fb,myspec,'r-',fb(indf)+[0 0],[err_low err_high].*myspec(indf),'k-');
set(h(end),'linewidth',2)

grid

%% label
ylabel('\Phi_{Vel} / (m/2)^2 / cpd')
xlabel('frequency / cpd')
```

indf =


502

ans =

1.5968e-05

err_high =

2.3163



It's nice to see that whatever our choice of M , our error bars agree fairly well with the observed spread. So our χ^2 model appears to be doing a decent job.

1.0.1 Getting the units right

We've talked about units for spectra, but let's lay everything out in one place. Our fundamental principle is that we want Parseval's theorem to work. But this gets a tiny bit messy when we average multiple frequencies. Still the basic rule of Parseval's theorem is not that the sum of the squares equals the sum of the squared Fourier coefficients, but rather that the integrated variance equals the integral under the spectrum.

Some instructions say to divide the spectral estimate by N^2 and then divide by $df = 1/T$, in effect multiplying by T/N^2 . But my sample code so far has only had a division by N . Why?

In the discrete Fourier transform, we have:

$$\sum_{n=0}^{N-1} x_n^2 = \frac{1}{N} \sum_{k=0}^{N-1} |X_k|^2,$$

or after spectral normalizations, maybe we write

$$\begin{aligned} \sum_{n=0}^{N-1} x_n^2 &= \frac{1}{N} \left[|X_0|^2 + \sum_{k=1}^{(N/2)-1} 2|X_k|^2 + |X_{N/2}|^2 \right] \\ &\approx \frac{1}{N} \sum_{k=0}^{(N-1)/2} 2|X_k|^2. \end{aligned}$$

For this discussion, to keep the equations compact, we'll use the approximation in the final line, neglecting the fact that the mean and the $k = N/2$ value should not be doubled. Many times we work with this form, since it gives us meaningful spectral slopes.

If I take the limiting case in which x_n has one frequency, then I can easily determine how to normalize X_0 . Suppose x_n is a constant, so $x_n = \bar{x}$. Then

$$\sum_{n=0}^{N-1} x_n^2 = N\bar{x}^2 = \frac{1}{N} \sum_{k=0}^{N-1} |X_k|^2 = \frac{1}{N} |X_0|^2,$$

since only frequency 0 has a non-zero Fourier coefficient. This implies that:

$$\bar{x} = \frac{|X_0|}{N}$$

Similarly, if $x = a \cos(\omega t)$, then

$$\sum_{n=0}^{N-1} x_n^2 = \sum_{n=0}^{N-1} a^2 \cos^2(\omega t) = \frac{Na^2}{2} = \frac{1}{N} \sum_{k=0}^{N-1} |X_k|^2 = \frac{2}{N} |X_0|^2,$$

which tells us how to determine the amplitude a :

$$a = \frac{2|X_0|}{N}.$$

In continuous form, we had a form more like this:

$$\int_{-\infty}^{\infty} x^2(t) dt = \int_{-\infty}^{\infty} |X(f)|^2 df$$

where f is frequency in cycles per unit time (or we might use $\sigma = 2\pi f$ sometimes, where σ is frequency in radians per unit time.) If we want this integral form to work for our real data, then we have to be a bit careful with our normalizations. We're going to want the area under the curve in our spectrum to be equal to the total variance integrated over time. So if total integrated variance is

$$variance = \sum_{n=0}^{N-1} x_n^2 \Delta t$$

where $\Delta t = T/N$. Then the integrated spectrum should be

$$variance = \frac{\alpha}{N} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f,$$

where $\Delta f = 1/(N\Delta t) = 1/T$, and we'll need to figure out α to ensure that the spectral estimator that we compute still properly adheres to Parseval's theorem. This implies that we might imagine normalizing our spectra to have:

$$\sum_{n=0}^{N-1} x_n^2 \Delta t = \frac{\alpha}{N} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f = \frac{\Delta t}{N\Delta f} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f = \frac{T^2}{N^2} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f$$

Or maybe this makes for units that aren't easily compared, so we could normalize our spectra to represent the average energy per unit time in the time domain, and adjust the frequency domain accordingly:

$$\begin{aligned} \frac{1}{T} \sum_{n=0}^{N-1} x_n^2 \Delta t &= \frac{1}{N\Delta t} \sum_{n=0}^{N-1} x_n^2 \Delta t \\ &= \frac{T}{N^2} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f \\ &= \frac{(\Delta t)}{N} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f \\ &= \frac{1}{N^2 \Delta f} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f \end{aligned}$$

This uses a normalization with an additional factor of $N \Delta f$ relative to what we got out of our initial discrete Fourier transform.

Alternatively, you could choose to express this in terms of the Nyquist frequency

$$\frac{1}{T} \sum_{n=0}^{N-1} x_n^2 \Delta t = \frac{1}{N^2 \Delta f} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f \tag{1}$$

$$= \frac{1}{N^2 f_{Nyquist}} \sum_{k=0}^{N/2-1} 2|X_k|^2 \Delta f \tag{2}$$

so we could divide our spectrum by twice the Nyquist frequency to have energy in units appropriate for comparing if we wanted to have our integrals match.

This isn't always the way we think about this, but it serves as our reminder that we should think about the units of our spectrum. What we know is that integral of our spectrum over a certain frequency range should give a measure of the signal variance:

$$\text{variance in a band} = \int_{f-\Delta f/2}^{f+\Delta f/2} |X(f)|^2 df \quad (3)$$

So if we expand this out, this implies that the units of $|X(f)|^2$ should be equivalent to variance divided by frequency, so it's our reminder that we'll label the y-axis units as the squared units of x divided by frequency, with a normalization to account for the units of time in our data.

The Fourier Transform of a Boxcar In the next lecture we'll look at "windowing" and "spectral leakage." To introduce these concepts, let's start here by reminding ourselves one immediate impact of finite records compared to infinite ones.

When we deal with data records of finite length, we're always taking finite sized segments of data, sort of like Fourier transforming a boxcar filter. Usually we think of a Fourier transform of this form:

$$X(f) = \int_{-\infty}^{\infty} x(t)e^{-i2\pi ft} dt.$$

If $x(t)$ is a boxcar filter, this becomes:

$$\int_{-\infty}^{\infty} h(t)e^{-i2\pi ft} dt = \int_{-1/2}^{1/2} e^{-i2\pi ft} dt = \left. \frac{e^{-i2\pi ft}}{-i2\pi f} \right|_{-1/2}^{1/2} = \frac{e^{i2\pi f/2} - e^{-i2\pi f/2}}{i2\pi f} = \frac{\sin(\pi f)}{\pi f}$$

This is the sinc function, which can be written as:

$$\text{sinc}(x) = \frac{\sin(x)}{x}$$

or (for digital signal processing)

$$\text{sinc}(x) = \frac{\sin(\pi x)}{\pi x}.$$

We'll try to use the second definition, the "normalized sinc function," since we're trying to use a frequency f that needs to be multiplied by 2π .

Remember that convolution in the time domain corresponds to multiplication in the frequency domain. In this case, we're doing the opposite. We multiplied our data by a boxcar window in the time domain, and that's equivalent to convolving with a sinc function in the frequency domain:

$$\hat{X}(f) = \int_{-\infty}^{\infty} X(g)W(f-g) dg$$

where

$$W(f) = \frac{\sin(\pi f T)}{\pi f T} = \text{sinc}(fT),$$


for a record length of T .

Correspondingly, the spectrum is convolved with $W(f)^2 = \text{sinc}(fT)^2$. This has a central peak of width $\Delta f = \pm 1/(T)$. The window will have a maximum value

$$|W(f)|^2 < (\pi f T)^{-2}$$

What does this actually look like and how bad is it for our spectral estimates?


```
[25]: T=30;
df=1/T;
fn=24;
f = linspace(-5,5)./T;
%nanmean(diff(f))
%df
f=-fn:(df/10):fn; %plot finer than our usual frequency res to see response.
y = sinc(f.*T).^2;
semilogy(f,y,f,1./(pi*f.*T).^2,'k--')
ylim([1e-5 1e0])
xlim([-5 5]/T);
%xlim([-1 1])
freqline(1/T);
freqline(-1/T);
ylabel('sinc^2(fT)')
xlabel('Frequency (cycles per day)')
```



So as long as our spectrum has similar energy at all frequencies (a “white” spectrum) the effects of the leakage will be limited, because the sidelobes decay as $1/f^2$. But if we have a “red” or a “blue” spectrum, we’ll have issues:

```
[26]: Strue=ones(size(f)); %a white "true" spectrum
Strue=Strue./f.^2;
Strue(find(f==0))=0;
%our response is sinc^2(fT).
%resp=sinc(f.*T).^2;
%Now our estimated spectrum is the convolution of the true spectrum times
%→sinc^2.
%df
Sest=conv(Strue,y,'same')./T;
semilogy(f,Strue,f,Sest(1:length(f)))
    xlim([-5 5]/T);

%plot(Sest./10)
%semilogy(f,resp);
%    xlim([-5 5]/T);
```



This is why you will sometimes hear of “pre whitening” spectra in order to minimize leakage. What we mean by this is that if we have a red f^{-2} spectrum, we can effectively multiply it by f^2 by differentiating our time series (recall the effect of differentiating on the Fourier coefficients from our previous lecture). If what we are Fourier transforming has a nearly white spectrum, leakage is minimized.

The side lobes of the window are definitely problematic. And it turns out that we’re not stuck with them. If we can widen the central peak of W (the Fourier transform of our window), we can lower the impact of the side lobes. To do this, we’ll want to forego the rectangular window in the time domain and choose something that lets us attenuate the beginnings and ends of each segment of our data. What if we chose a triangle window? That will already give us fewer side lobes.

But we can keep going to find a window that looks more like a Gaussian. Leading possibilities:

Cosine taper:


$$w(t) = \cos^\alpha \left(\frac{\pi t}{T} \right)$$

with $\alpha = [1, 4]$.

Hanning window or “raised cosine” window (developed by von Hann):

$$w(t) = \cos^2\left(\frac{\pi t}{T}\right) = \frac{1 + \cos(2\pi t/T)}{2} = 0.5 + 0.5 \cos(2\pi t/T)$$

```
[138]: time=-0.5:0.01:0.5;
T=1;
hanwin=cos(pi*time / T).^2;
plot(time,hanwin,time,cos(pi*time/(T)))
legend('Hanning','cosine')
```



```
[ ]:
```